

Assignment 2:

[IBM-Project-117-1658211897](#)

Importing

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings("ignore")
```

In [1]:

```
df=pd.read_csv("E:\IBM projects Assignment Sona
College\Churn_Modelling.csv")
```

In [4]:

```
df.info()

RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   RowNumber             10000 non-null  int64  
 1   CustomerId            10000 non-null  int64  
 2   Surname               10000 non-null  object  
 3   CreditScore           10000 non-null  int64  
 4   Geography             10000 non-null  object  
 5   Gender               10000 non-null  object  
 6   Age                  10000 non-null  int64  
 7   Tenure               10000 non-null  int64  
 8   Balance              10000 non-null  float64 
 9   NumOfProducts        10000 non-null  int64  
10   HasCrCard            10000 non-null  int64  
11   IsActiveMember       10000 non-null  int64  
12   EstimatedSalary      10000 non-null  float64 
13   Exited               10000 non-null  int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In [5]:

```
df.describe()
```

In [6]:

```
Out[6]:
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000	10000	10000	10000	10000	10000	10000	10000	10000	10000	10000
unique	10000	10000	10000	10000	10000	10000	10000	10000	10000	10000	10000
top	0	4	0	0	0	000000	00000	0	00000	00000	0

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCreditCard	IsActiveMember	EstimatedSalary	Exited
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.70550	0.515100	100090.239881	0.203700
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.45584	0.499797	57510.492818	0.402769
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.000000	0.000000	11.580000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.000000	0.000000	51002.110000	0.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.000000	1.000000	100193.915000	0.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.000000	1.000000	149388.247500	0.000000
max	10000.000000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.000000	1.000000	199992.480000	1.000000

1. UNIVARIATE ANALYSIS


The term univariate analysis refers to the analysis of one variable. You can remember this because the prefix “uni” means “one.” There are three common ways to perform univariate analysis on one variable: 1. Summary statistics – Measures the center and spread of values.

In [78]:

```
df_france=df.loc[df['Geography']=='France']
df_spain=df.loc[df['Geography']=='Spain']
df_germany=df.loc[df['Geography']=='Germany']
```

In [79]:

```
plt.plot(df_france['Balance'],np.zeros_like(df_france['Balance']),'o')
plt.plot(df_spain['Balance'],np.zeros_like(df_spain['Balance']),'o')
plt.plot(df_germany['Balance'],np.zeros_like(df_germany['Balance']),'o')
plt.xlabel('Age')
plt.show()
```




Histogram

```
sns.displot(df["Age"], color='darkorange')
```

In [7]:

Out[7]:




In Histogram, we can do it vertically too, by just changing the axis

```
sns.histplot(y="Age", data=df, color='darkorange')
```

In [8]:

Out[8]:




Now, we can also use Histogram for categorical variables

```
sns.histplot(x='Age',data=df,hue=df['Tenure'])
```

In [9]:

Out[9]:




Distplot

```
sns.distplot(df["Age"],color='purple')
```

In [10]:

Out[10]:




This is visualising distplot alone

In [11]:

```
sns.distplot(df["Age"], hist=False, color='blue')
```

Out[11]:




Boxplot

In [12]:

```
sns.boxplot(df["Age"], color='pink')
```

Out[12]:




Countplot

In [14]:

```
sns.countplot(df['Age'])
```

Out[14]:




2. BIVARIATE ANALYSIS

Image result for bivariate analysis in python It is a methodical statistical technique applied to a pair of variables (features/ attributes) of data to determine the empirical relationship between them. In order words, it is meant to determine any concurrent relations (usually over and above a simple correlation analysis).

In [80]:

```
sns.FacetGrid(df,hue="Geography",size=5).map(plt.scatter,"Age","Balance").a  
dd_legend();  
plt.show()
```

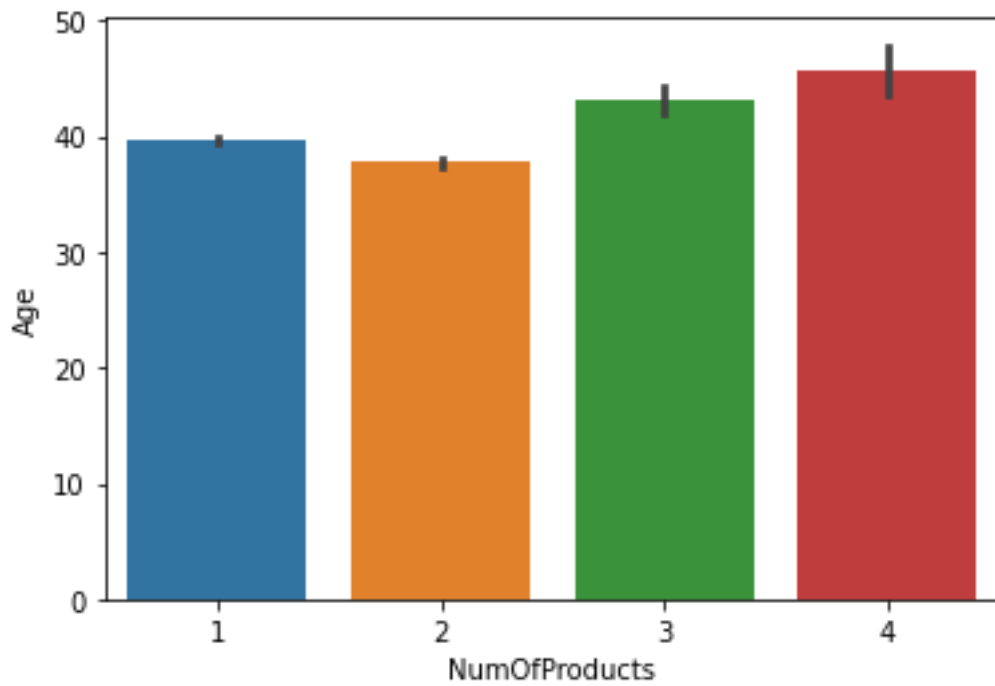


Barplot

In [15]:

```
sns.barplot(df["NumOfProducts"],df["Age"])
```


Out[15]:



Linearplot

In [16]: `sns.lineplot(df["Age"],df["NumOfProducts"], color='purple')`


Out[16]:



Scatterplot

In [17]: `sns.scatterplot(x=df.Age,y=df.RowNumber,color='green')`

Out[17]:




Pointplot

In [18]:

```
sns.pointplot(x='Age',y='Tenure',data=df,color='pink')
```

Out[18]:




Regplot

In [19]:

```
sns.regplot(df['Age'],df['Tenure'],color='orange')
```

Out[19]:



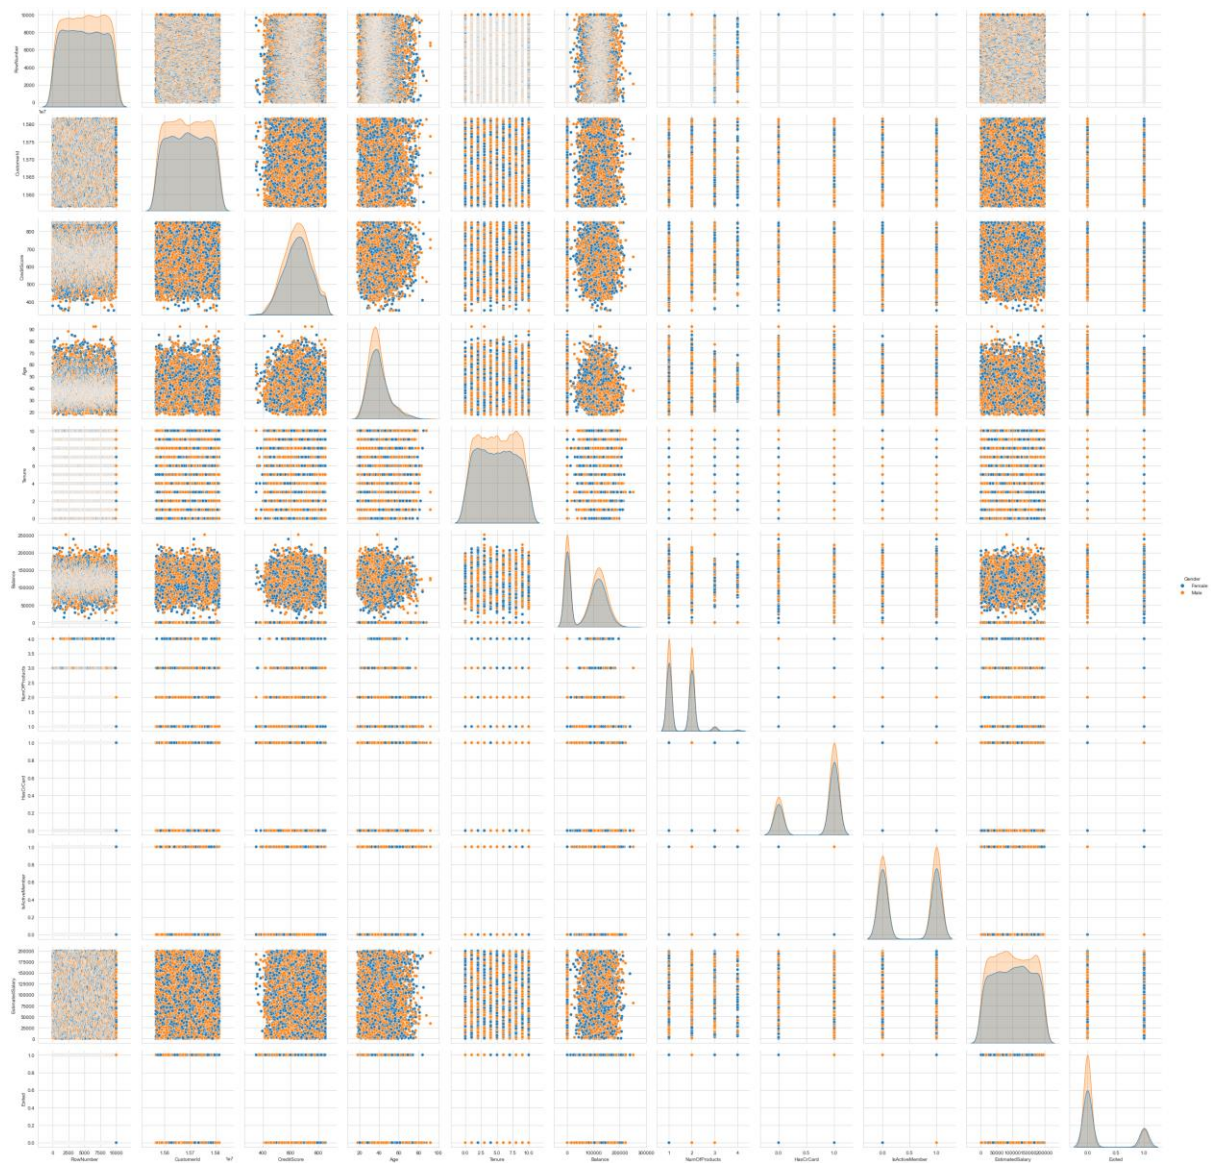
3. MULTI - VARIATE ANALYSIS

Multivariate analysis is based in observation and analysis of more than one statistical outcome variable at a time. In design and analysis, the technique is used to perform trade studies across multiple dimensions while taking into account the effects of all variables on the responses of interest.

In [187]:

```
sns.pairplot(df,hue="Gender",size=3)
```

Out[187]:

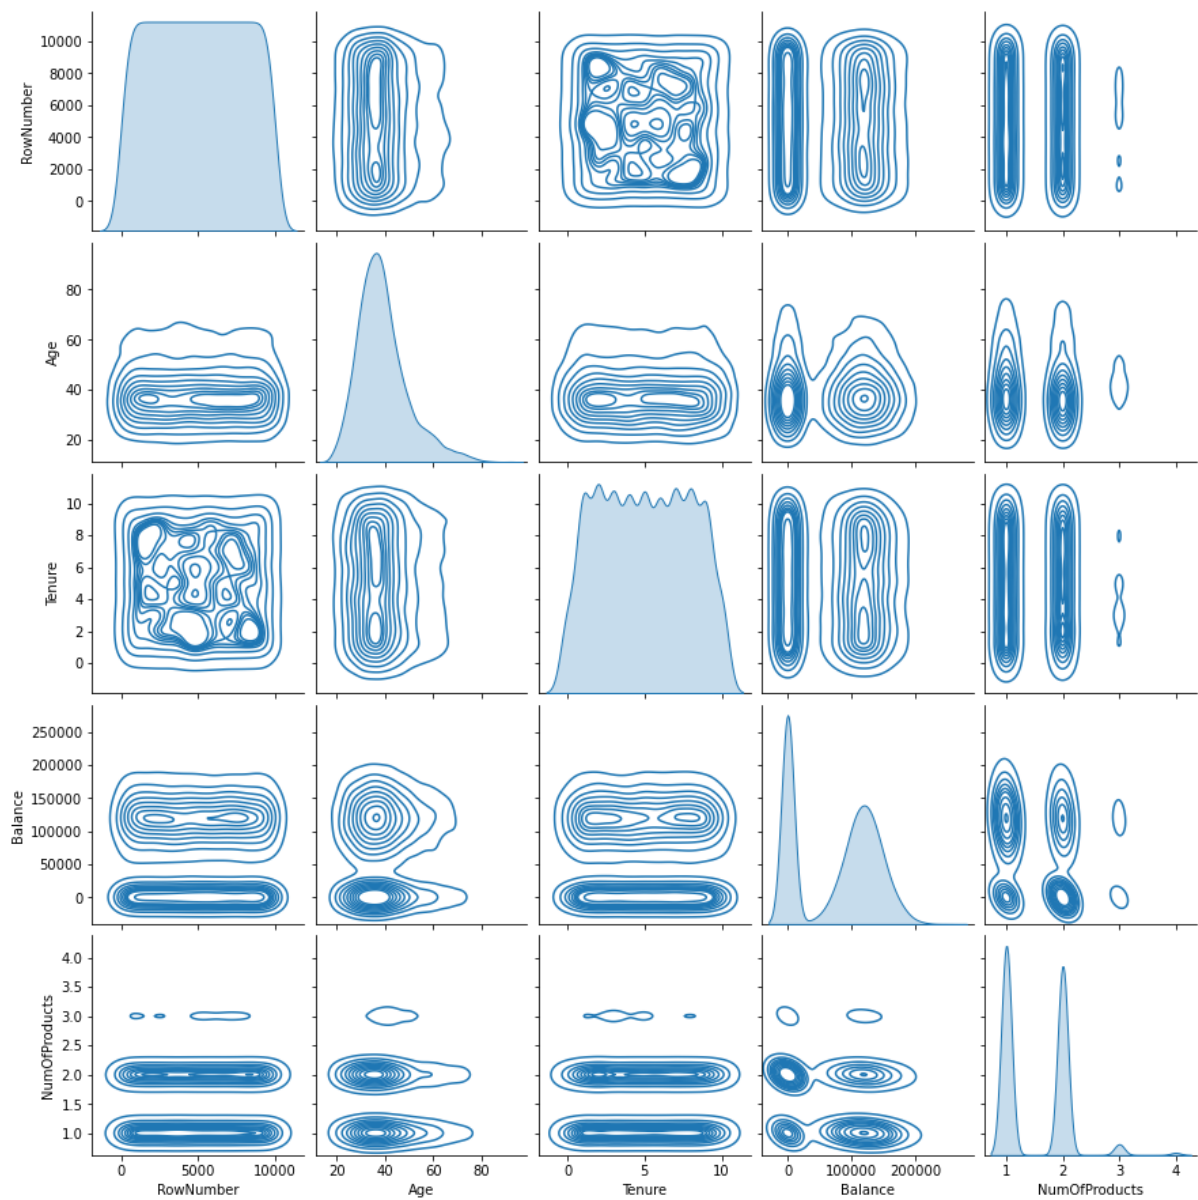


Pairplot

In [21]:

```
sns.pairplot(data=df[["RowNumber", "Age", "Tenure", "Balance", "NumOfProducts"]], kind="kde")
```

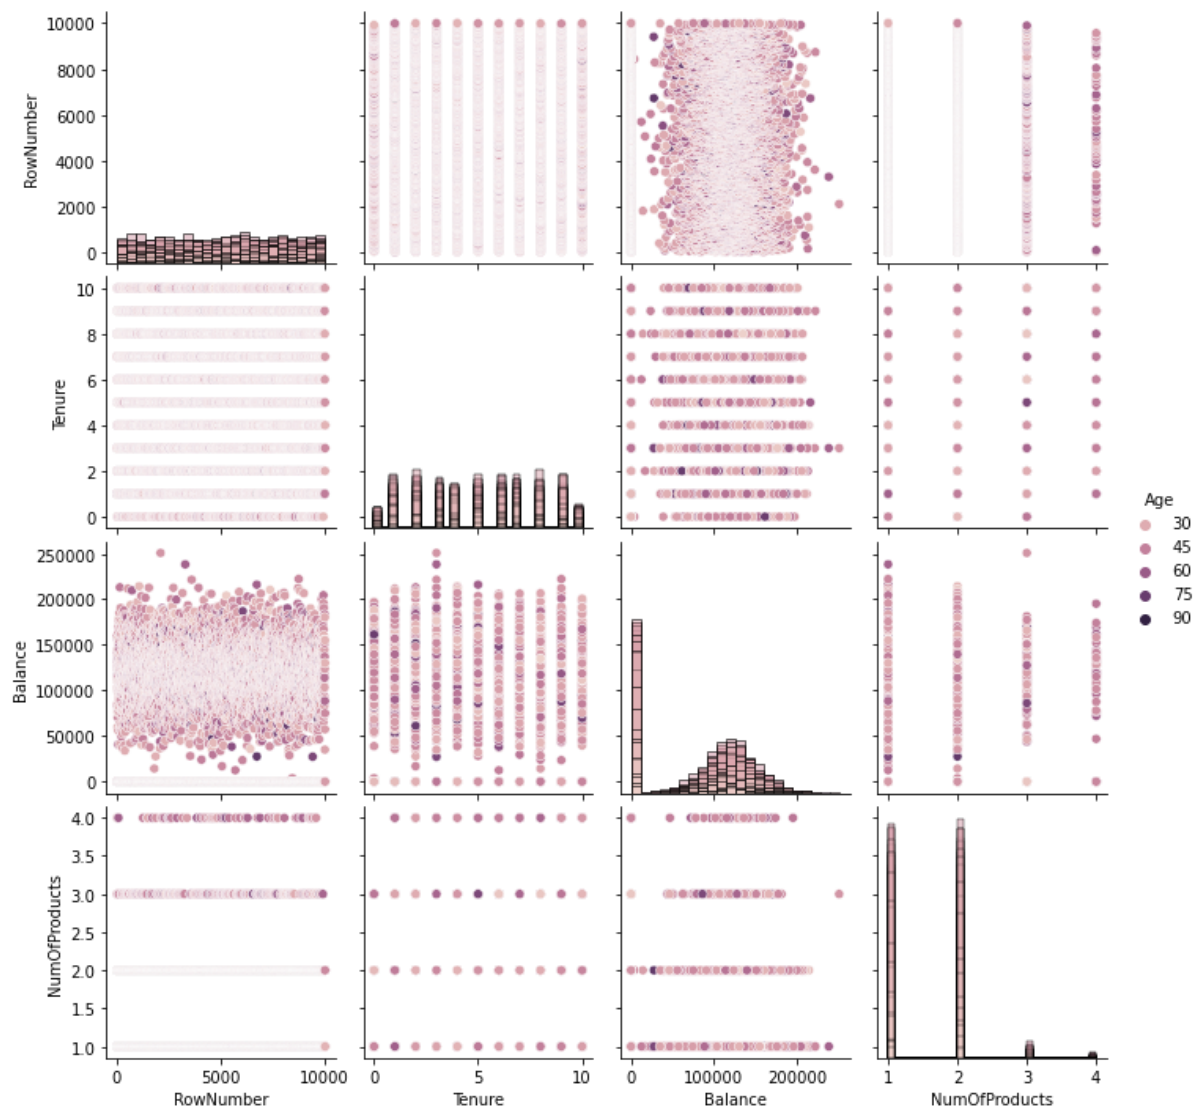
Out[21]:



In [22]:

```
sns.pairplot(data=df[["RowNumber", "Age", "Tenure", "Balance", "NumOfProducts"]], hue="Age", diag_kind="hist")
```

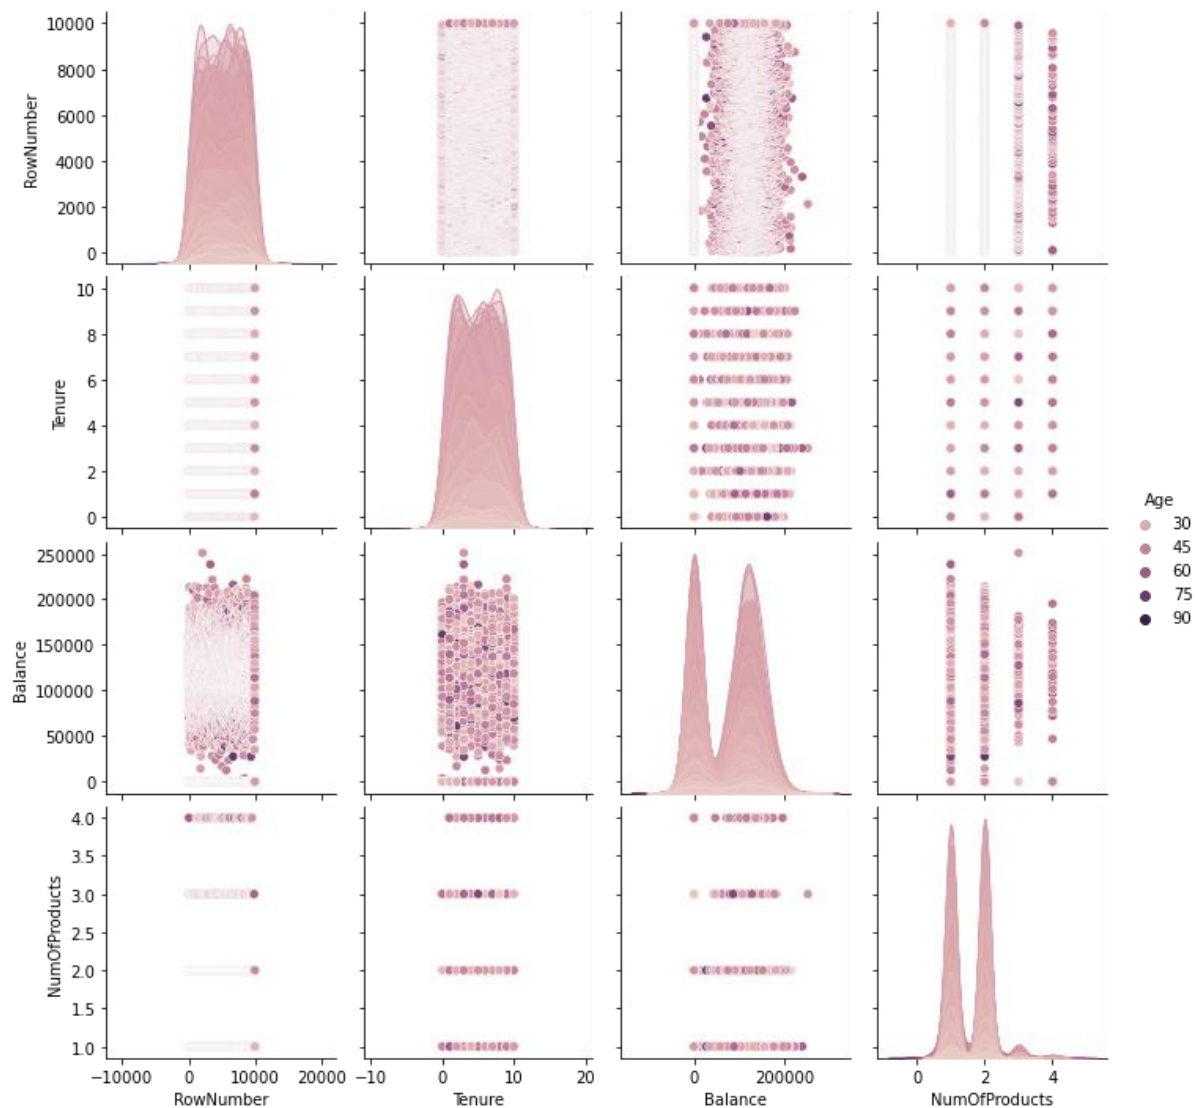
Out[22]:



In [23]:

```
sns.pairplot(data=df[["RowNumber", "Age", "Tenure", "Balance", "NumOfProducts"]], hue="Age")
```

Out[23]:



4. Perform descriptive statistics on the dataset

Image result for descriptive statistics in python Python Descriptive Statistics process describes the basic features of data in a study. It delivers summaries on the sample and the measures and does not use the data to learn about the population it represents. Under descriptive statistics, fall two sets of properties- central tendency and dispersion.

In [24]:

```
df.describe()
```

Out[24]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
count	10000.000000	1.00000e+04	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
mean	5000.500000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	1.530200	0.705500	0.515100	100090.239881	0.203700
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	0.581654	0.455840	0.499797	57510.492818	0.402769
min	1.000000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	1.000000	0.000000	0.000000	11.580000	0.000000
25%	2500.750000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	1.000000	0.000000	0.000000	51002.100000	0.000000
50%	5000.500000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	1.000000	1.000000	1.000000	100193.915000	0.000000
75%	7500.250000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	2.000000	1.000000	1.000000	149388.247500	0.000000
max	10000.000000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	4.000000	1.000000	1.000000	199992.480000	1.000000

In [81]:

df.head()

Out[81]:

	Row Number	Customer Id	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	1	0	42	2	0.00	1	1	1	101348.88	1

	Row Number	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
1	2	15647311	Hill	608	2	0	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	1	0	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	1	0	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	2	0	43	2	125510.82	1	1	1	79084.10	0

In [82]:

```
df.mean() # Get the mean of each column
```

Out[82]:

```
RowNumber      5.000500e+03
CustomerId      1.569094e+07
CreditScore     6.505288e+02
Geography       1.749500e+00
Gender          5.457000e-01
Age            4.019530e+01
Tenure          5.012800e+00
Balance         7.648589e+04
NumOfProducts   1.530200e+00
HasCrCard       7.055000e-01
IsActiveMember  5.151000e-01
EstimatedSalary 1.000902e+05
Exited          2.037000e-01
dtype: float64
```

In [83]:

```
df.mean(axis=1) # Get the mean of each row
```

Out[83]:

```
0      1.210509e+06
1      1.218794e+06
2      1.222574e+06
3      1.215071e+06
4      1.226414e+06
...
9995   1.208717e+06
9996   1.210733e+06
9997   1.202875e+06
9998   1.220088e+06
9999   1.215960e+06
Length: 10000, dtype: float64
```

In [84]:

```
df.median()                                # Get the median of each column
```

Out[84]:


```
RowNumber      5.000500e+03
CustomerId     1.569074e+07
CreditScore    6.520000e+02
Geography      1.000000e+00
Gender         1.000000e+00
Age            3.800000e+01
Tenure         5.000000e+00
Balance        9.719854e+04
NumOfProducts 1.000000e+00
HasCrCard      1.000000e+00
IsActiveMember 1.000000e+00
EstimatedSalary 1.001939e+05
Exited         0.000000e+00
dtype: float64
```

In [85]:

```
norm_data = pd.DataFrame(np.random.normal(size=100000))

norm_data.plot(kind="density",
               figsize=(10,10));
plt.vlines(norm_data.mean(),      # Plot black line at mean
           ymin=0,
           ymax=0.4,
           linewidth=5.0);

plt.vlines(norm_data.median(),    # Plot red line at median
           ymin=0,
           ymax=0.4,
           linewidth=2.0,
           color="red");
```




In [86]:

```
skewed_data = pd.DataFrame(np.random.exponential(size=100000))

skewed_data.plot(kind="density",
                  figsize=(10,10),
                  xlim=(-1,5));

plt.vlines(skewed_data.mean(),      # Plot black line at mean
            ymin=0,
            ymax=0.8,
            linewidth=5.0);

plt.vlines(skewed_data.median(),    # Plot red line at median
            ymin=0,
            ymax=0.8,
            linewidth=2.0,
            color="red");
```



In [87]:

```
norm_data = np.random.normal(size=50)
outliers = np.random.normal(15, size=3)
combined_data = pd.DataFrame(np.concatenate((norm_data, outliers), axis=0))

combined_data.plot(kind="density",
                    figsize=(10,10),
                    xlim=(-5,20));

plt.vlines(combined_data.mean(),      # Plot black line at mean
            ymin=0,
            ymax=0.2,
            linewidth=5.0);

plt.vlines(combined_data.median(),   # Plot red line at median
            ymin=0,
            ymax=0.2,
            linewidth=2.0,
            color="red");
```



```
df.mode()
```

[illegible]

	Row Num ber	Cust omer Id	Sur na me	Cred itSco re	Geo grap hy	Ge nd er	A g e	Te nu re	Bal anc e	NumO fProdu cts	Has CrC ard	IsActiv eMem ber	Estima tedSala ry	Ex ite d
3	4	1556 5779	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	5	1556 5796	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
...
9995	9996	1581 5628	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9996	9997	1581 5645	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9997	9998	1581 5656	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9998	9999	1581 5660	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
9999	10000	1581 5690	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

10000 rows × 14 columns

Measures of Spread

In [89]:

```
max(df["Age"]) - min(df["Age"])
```

Out[89]:

67

In [90]:

```
five_num = [df["Age"].quantile(0),
             df["Age"].quantile(0.25),
             df["Age"].quantile(0.50),
```

```
df["Age"].quantile(0.75),  
df["Age"].quantile(1)]
```

```
five_num
```

```
[25.0, 33.0, 38.0, 46.0, 92.0]
```

```
df["Age"].describe()
```

```
count      10000.000000  
mean         40.195300  
std         10.047729  
min          25.000000  
25%          33.000000  
50%          38.000000  
75%          46.000000  
max          92.000000  
Name: Age, dtype: float64
```

```
df["Age"].quantile(0.75) - df["Age"].quantile(0.25)
```

```
13.0
```

```
df.boxplot(column="Age",  
            return_type='axes',  
            figsize=(8,8))
```

```
plt.text(x=0.74, y=22.25, s="3rd Quartile")  
plt.text(x=0.8, y=18.75, s="Median")  
plt.text(x=0.75, y=15.5, s="1st Quartile")  
plt.text(x=0.9, y=10, s="Min")  
plt.text(x=0.9, y=33.5, s="Max")  
plt.text(x=0.7, y=19.5, s="IQR", rotation=90, size=25);
```

Out[90]:


In [91]:

Out[91]:

In [92]:

Out[92]:

In [93]:



<code>df["Age"].var()</code>	In [94]:
100.95685359535904	Out[94]:
<code>df["Age"].std()</code>	In [95]:
10.047728777955694	Out[95]:
<code>abs_median_devs = abs(df["Age"] - df["Age"].median())</code>	In [96]:
<code>abs_median_devs.median() * 1.4826</code>	Out[96]:
8.8956	

Skewness and Kurtosis

```
df["Age"].skew()    # Check skewness
```

In [97]:

```
1.0495460120728233
```

Out[97]:

```
df["Age"].kurt()    # Check kurtosis
```

In [98]:

```
1.2747003702904487
```

Out[98]:


In [99]:

```
norm_data = np.random.normal(size=100000)
skewed_data = np.concatenate((np.random.normal(size=35000)+2,
                               np.random.exponential(size=65000)),
                              axis=0)
uniform_data = np.random.uniform(0,2, size=100000)
peaked_data = np.concatenate((np.random.exponential(size=50000),
                               np.random.exponential(size=50000)*(-1)),
                              axis=0)

data_df = pd.DataFrame({"norm":norm_data,
                        "skewed":skewed_data,
                        "uniform":uniform_data,
                        "peaked":peaked_data})

data_df.plot(kind="density",
              figsize=(10,10),
              xlim=(-5,5));
```

In [100]:



In [101]:

```
data_df.skew()
```

Out[101]:

```
norm      0.009355
skewed    1.019859
uniform   -0.003685
peaked    -0.010567
dtype: float64
```

In [102]:

```
data_df.kurt()
```

Out[102]:

```
norm      0.019495
skewed    1.526498
uniform   -1.201582
peaked    3.080928
dtype: float64
```

5. Handle the Missing values.

```
df.head()
```

In [103]:

Out[103]:

	Row Num ber	Cust omer Id	Sur na me	Cred itSco re	Geo grap hy	Ge nd er	A g e	Te nu re	Bal anc e	NumO fProdu cts	Has CrC ard	IsActiv eMem ber	Estima tedSala ry	Ex ite d
0	1	1563 4602	Har gra ve	619	1	0	4 2	2	0.00	1	1	1	101348 .88	1
1	2	1564 7311	Hill	608	2	0	4 1	1	838 07.8 6	1	0	1	112542 .58	0
2	3	1561 9304	Oni o	502	1	0	4 2	8	159 660. 80	3	1	0	113931 .57	1
3	4	1570 1354	Bon i	699	1	0	3 9	1	0.00	2	0	0	93826. 63	0
4	5	1573 7888	Mit chel l	850	2	0	4 3	2	125 510. 82	1	1	1	79084. 10	0

```
df.isnull()
```

In [104]:

Out[104]:


	Row Num ber	Cust omer Id	Sur na me	Cred itSco re	Geo grap hy	Ge nd er	A g e	Te nu re	Bal anc e	NumO fProdu cts	Has CrC ard	IsActiv eMem ber	Estima tedSala ry	Ex ite d
0	False	False	Fals e	False	False	Fals e	F al se	Fals e	Fals e	False	False	False	False	Fa lse
1	False	False	Fals e	False	False	Fals e	F al se	Fals e	Fals e	False	False	False	False	Fa lse
2	False	False	Fals e	False	False	Fals e	F al se	Fals e	Fals e	False	False	False	False	Fa lse
3	False	False	Fals e	False	False	Fals e	F al se	Fals e	Fals e	False	False	False	False	Fa lse

	Row Num ber	Cust omer Id	Sur na me	Cred itSco re	Geo grap hy	Ge nder	A ge	Te nu re	Bal anc e	NumO fProdu cts	Has CrC ard	IsActiv eMem ber	Estima tedSala ry	Ex ited
4	False	False	False	False	False	False	False	False	False	False	False	False	False	False
...
9995	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9996	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9997	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9998	False	False	False	False	False	False	False	False	False	False	False	False	False	False
9999	False	False	False	False	False	False	False	False	False	False	False	False	False	False

10000 rows × 14 columns

```
In [105]: sns.heatmap(df.isnull(),yticklabels=False,cbar=False,cmap='viridis')
```


Out[105]:



In [106]:

```
sns.set_style('whitegrid')
sns.countplot(x='Geography', data=df)
```


Out[106]:



In [107]:

```
sns.set_style('whitegrid')
sns.countplot(x='Geography', hue='Gender', data=df, palette='RdBu_r')
```


Out[107]:



In [108]:

```
sns.set_style('whitegrid')  
sns.countplot(x='Geography',hue='Gender',data=df,palette='rainbow')
```


Out[108]:



In [109]:

```
sns.distplot(df['Age'].dropna(),kde=False,color='darkred',bins=40)
```


Out[109]:



In [110]:

```
df['Age'].hist(bins=30,color='darkred',alpha=0.3)
```


Out[110]:



In [111]:

```
sns.countplot(x='NumOfProducts',data=df)
```


Out[111]:



In [112]:

```
df['Age'].hist(color='green',bins=40,figsize=(8,4))
```

Out[112]:



Cufflinks for plots

In [113]:

```
import cufflinks as cf
cf.go_offline()
```

ModuleNotFoundError

Traceback (most recent call last)

Input In [113], in ()

```
----> 1 import cufflinks as cf
      2 cf.go_offline()
```


ModuleNotFoundError: No module named 'cufflinks'

In [114]:

```
df['Age'].iplot(kind='hist',bins=30,color='green')
```

AttributeError Traceback (most recent call last)

Input In [114], in ()

```
----> 1 df['Age'].iplot(kind='hist',bins=30,color='green')
```

File ~\anaconda3\lib\site-packages\pandas\core\generic.py:5575, in NDFrame.
__getattr__(self, name)

```
5568 if (
5569     name not in self._internal_names_set
5570     and name not in self._metadata
5571     and name not in self._accessors
5572     and self._info_axis._can_hold_identifiers_and_holds_name(name)
5573 ):
5574     return self[name]
-> 5575 return object.__getattribute__(self, name)
```


AttributeError: 'Series' object has no attribute 'iplot'

Data Cleaning

In [115]:

```
plt.figure(figsize=(12, 7))
sns.boxplot(x='Gender',y='Age',data=df,palette='winter')
```

Out[115]:



In [116]:

```
def impute_age(cols):
    Age = cols[0]
    Pclass = cols[1]

    if pd.isnull(Age):
```

```

    if Pclass == 1:
        return 37

    elif Pclass == 2:
        return 29


    else:
        return 24

else:
    return Age

```

In [117]:
sns.heatmap(df.isnull(),yticklabels=False,cbar=False,cmap='viridis')

Out[117]:



In [118]:
df.drop('Gender',axis=1,inplace=True)

In [119]:
df.head()

Out[119]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	1	42	2	0.00	1	1	1	101348.88	1

	RowNumber	CustomerId	Surname	CreditScore	Geography	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
1	2	15647311	Hill	608	2	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	Onio	502	1	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	1	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Mitchell	850	2	43	2	125510.82	1	1	1	79084.10	0

Converting Categorical Features

In [120]:

```
df.info()
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   RowNumber              10000 non-null  int64  
1   CustomerId             10000 non-null  int64  
2   Surname                10000 non-null  object  
3   CreditScore             10000 non-null  int64  
4   Geography               10000 non-null  int64  
5   Age                    10000 non-null  int64  
6   Tenure                  10000 non-null  int64  
7   Balance                 10000 non-null  float64 
8   NumOfProducts          10000 non-null  int64  
9   HasCrCard              10000 non-null  int64  
10  IsActiveMember         10000 non-null  int64  
11  EstimatedSalary         10000 non-null  float64 
12  Exited                  10000 non-null  int64  
dtypes: float64(2), int64(10), object(1)
memory usage: 1015.8+ KB
```

In [121]:

```
pd.get_dummies(df['Geography'],drop_first=True).head()
```

Out[121]:

```

2  3
0  0  0
```

```
2 3
1 1 0
2 0 0
3 0 0
4 1 0
```

```
df.info
```

In [122]:

Out[122]:

In [123]:

```
sex = pd.get_dummies(df['Age'],drop_first=True)
embark = pd.get_dummies(df['Balance'],drop_first=True)
```

In [124]:

```
df.drop(['Age','HasCrCard','Surname','CustomerId'],axis=1,inplace=True)
```

In [125]:

```
df.head()
```

Out[125]:

	RowNum ber	CreditSc ore	Geogra phy	Tenu re	Balanc e	NumOfProd ucts	IsActiveMe mber	EstimatedSa lary	Exit ed
0	1	619	1	2	0.00	1	1	101348.88	1
1	2	608	2	1	83807.8 6	1	1	112542.58	0
2	3	502	1	8	159660. 80	3	0	113931.57	1
3	4	699	1	1	0.00	2	0	93826.63	0
4	5	850	2	2	125510. 82	1	1	79084.10	0

In [126]:

```
train = pd.concat([df,sex,embark],axis=1)
```

In [127]:

```
train.head()
```

Out[127]:

Row Number	Credit Score	Geography	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited	2	.	2	1	2	2	2	2	2	2	2	2
0	1	619	1	2	0	1	1	101348.88	1	0	.	0	0	0	0	0	0	0	0	0
1	2	608	2	1	8	1	1	112542.58	0	0	.	0	0	0	0	0	0	0	0	0
2	3	502	1	8	9	3	0	113931.57	1	0	.	0	0	0	0	0	0	0	0	0
3	4	699	1	1	0	2	0	93826.63	0	0	.	0	0	0	0	0	0	0	0	0
4	5	850	2	2	5	1	1	79084.10	0	0	.	0	0	0	0	0	0	0	0	0

5 rows × 6452 columns

```
In [128]: data=pd.DataFrame({"a": [1,2,np.nan], "b": [1,np.nan,np.nan], "c": [1,2,4] })
data
```

Out[128]:

	a	b	c
0	1.0	1.0	1

	a	b	c
1	2.0	NaN	2
2	NaN	NaN	4

```
data.isnull().any()
```

In [129]:

```
a      True
b      True
c     False
dtype: bool
```

Out[129]:

```
data.isnull().sum()
```

In [130]:

```
a      1
b      2
c      0
dtype: int64
```

Out[130]:

```
data.fillna(value = "S")
```

In [131]:

	a	b	c
0	1.0	1.0	1
1	2.0	S	2
2	S	S	4

Out[131]:

```
data["a"].mean()
```

In [132]:

```
1.5
```

Out[132]:

```
data["a"].median()
```

In [133]:

```
1.5
```

Out[133]:

6. Find the outliers and replace the outliers

For data that follows a normal distribution, the values that fall more than three standard deviations from the mean are typically considered outliers. Outliers can find their way into a dataset naturally

through variability, or they can be the result of issues like human error, faulty equipment, or poor sampling.

In [68]:

```
dataset= [11,10,12,14,12,15,14,13,15,102,12,14,17,19,107,
10,13,12,14,12,108,12,11,14,13,15,10,15,12,10,14,13,15,10]
```

Detecting outlier using Z score

Using Z score

In [69]:

```
outliers=[]
def detect_outliers(data):

    threshold=3
    mean = np.mean(data)
    std =np.std(data)

    for i in data:
        z_score= (i - mean)/std
        if np.abs(z_score) > threshold:
            outliers.append(y)
    return outliers
```

In [70]:

```
outlier_pt=detect_outliers(dataset)
```

In [71]:

```
outlier_pt
```

Out[71]:

```
[0      Hargrave
 1      Hill
 2      Onio
 3      Boni
 4      Mitchell
 ...
9995   Obijiaku
9996   Johnstone
9997      Liu
9998   Sabbatini
9999      Walker
Name: Surname, Length: 10000, dtype: object,
0      Hargrave
1      Hill
2      Onio
3      Boni
4      Mitchell
 ...
9995   Obijiaku
9996   Johnstone
9997      Liu
9998   Sabbatini
```

```

9999      Walker
Name: Surname, Length: 10000, dtype: object,
0      Hargrave
1      Hill
2      Onio
3      Boni
4      Mitchell
...
9995      Obijiaku
9996      Johnstone
9997      Liu
9998      Sabbatini
9999      Walker
Name: Surname, Length: 10000, dtype: object]

```

In [72]:

```

## Perform all the steps of IQR
sorted(dataset)

```

Out[72]:

```

[10,
 10,
 10,
 10,
 10,
 11,
 11,
 12,
 12,
 12,
 12,
 12,
 12,
 12,
 12,
 13,
 13,
 13,
 13,
 13,
 14,
 14,
 14,
 14,
 14,
 14,
 15,
 15,
 15,
 15,
 15,
 17,
 19,
 102,
 107,
 108]

```

In [73]:

```

quantile1, quantile3= np.percentile(dataset, [25,75])

```

In [74]:


```
print(quantile1,quantile3)
12.0 15.0
```

In [75]:

```
## Find the IQR
```

```
iqr_value=quantile3-quantile1
print(iqr_value)
3.0
```

In [76]:

```
## Find the lower bound value and the higher bound value
```

```
lower_bound_val = quantile1 -(1.5 * iqr_value)
upper_bound_val = quantile3 +(1.5 * iqr_value)
```


In [77]:

```
print(lower_bound_val,upper_bound_val)
7.5 19.5
```

In [46]:

```
sns.boxplot(df["Age"],color='purple')
```

Out[46]:




In [47]:

```
df["Age"]=np.where(df["Age"]<25,50,df["Age"])
```

In [48]:

```
sns.boxplot(df["Age"],color='pink')
```

Out[48]:



7. Check for Categorical columns and perform encoding.

Categorical Columns : Categorical are a Pandas data type. A string variable consisting of only a few different values.

Encoding : For efficient storage of these strings, the sequence of code points is converted into a set of bytes. The process is known as encoding.

In [139]:

```
df=pd.read_csv("E:\IBM projects Assignment Sona
College\Churn_Modelling.csv")
```

In [140]:

```
df.head()
```

Out[140]:

	Row Number	Customer Id	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	Hargrave	619	France	Female	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	Hill	608	Spain	Female	41	1	83807.86	1	0	1	112542.58	0

	Row Number	Customer Id	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
2	3	15619304	Onio	502	France	Female	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	Boni	699	France	Female	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	Michell	850	Spain	Female	43	2	125510.82	1	1	1	79084.10	0

In [141]:

```
df_numeric = df[['RowNumber', 'CustomerId', 'CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited']]
df_categorical = df[['Surname', 'Geography', 'Gender']]
```

In [142]:

```
df_numeric.head()
```

Out[142]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	15634602	619	42	2	0.00	1	1	1	101348.88	1
1	2	15647311	608	41	1	83807.86	1	0	1	112542.58	0
2	3	15619304	502	42	8	159660.80	3	1	0	113931.57	1
3	4	15701354	699	39	1	0.00	2	0	0	93826.63	0
4	5	15737888	850	43	2	125510.82	1	1	1	79084.10	0

In [143]:

```
df_categorical.head()
```

Out[143]:

	Surname	Geography	Gender
0	Hargrave	France	Female
1	Hill	Spain	Female
2	Onio	France	Female
3	Boni	France	Female
4	Mitchell	Spain	Female

```
print(df['Surname'].unique())
print(df['Geography'].unique())
print(df['Gender'].unique())

['Hargrave' 'Hill' 'Onio' ... 'Kashiwagi' 'Aldridge' 'Burbidge']
['France' 'Spain' 'Germany']
['Female' 'Male']
```

```
from sklearn.preprocessing import LabelEncoder
```

```
marry_encoder = LabelEncoder()
```

```
marry_encoder.fit(df_categorical['Gender'])
```

```
LabelEncoder()
```

```
marry_values = marry_encoder.transform(df_categorical['Gender'])
```

```
print("Before Encoding:", list(df_categorical['Gender'][-10:]))
print("After Encoding:", marry_values[-10:])
print("The inverse from the encoding result:",
marry_encoder.inverse_transform(marry_values[-10:]))
```

```
Before Encoding: ['Male', 'Female', 'Male', 'Male', 'Female', 'Male', 'Male', 'Male', 'Female', 'Male']
After Encoding: [1 0 1 1 0 1 1 0 1 0]
The inverse from the encoding result: ['Male' 'Female' 'Male' 'Male' 'Female' 'Male' 'Male' 'Male' 'Female' 'Male']
```

```
residence_encoder = LabelEncoder()
residence_values =
residence_encoder.fit_transform(df_categorical['Geography'])
```

```
print("Before Encoding:", list(df_categorical['Geography'][:5]))
print("After Encoding:", residence_values[:5])
print("The inverse from the encoding result:",
residence_encoder.inverse_transform(residence_values[:5]))
```

```
Before Encoding: ['France', 'Spain', 'France', 'France', 'Spain']
After Encoding: [0 2 0 0 2]
The inverse from the encoding result: ['France' 'Spain' 'France' 'France' 'Spain']
```

In [150]:

```
from sklearn.preprocessing import OneHotEncoder
```

```
gender_encoder = OneHotEncoder()
```

In [151]:

```
from sklearn.preprocessing import OneHotEncoder
import numpy as np
```

```
gender_encoder = OneHotEncoder()
gender_resaped = np.array(df_categorical['Gender']).reshape(-1, 1)
gender_values = gender_encoder.fit_transform(gender_resaped)
```

```
print(df_categorical['Gender'][:5])
print()
print(gender_values.toarray()[:5])
print()
print(gender_encoder.inverse_transform(gender_values)[:5])
```

```
0    Female
1    Female
2    Female
3    Female
4    Female
Name: Gender, dtype: object
```

```
[[1. 0.]
 [1. 0.]
 [1. 0.]
 [1. 0.]
 [1. 0.]]
```

```
['Female']
['Female']
['Female']
['Female']
['Female']
```

In [152]:

```
smoke_encoder = OneHotEncoder()
smoke_resaped = np.array(df_categorical['Surname']).reshape(-1, 1)
smoke_values = smoke_encoder.fit_transform(smoke_resaped)
```

```
print(df_categorical['Surname'][:5])
print()
print(smoke_values.toarray()[:5])
print()
print(smoke_encoder.inverse_transform(smoke_values)[:5])
```

```
0    Hargrave
1      Hill
2      Onio
3      Boni
4    Mitchell
Name: Surname, dtype: object
```

```
[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]]
```

```
[['Hargrave']
 ['Hill']
 ['Onio']
 ['Boni']
 ['Mitchell']]
```

In [153]:

```
work_encoder = OneHotEncoder()
work_resaped = np.array(df_categorical['Geography']).reshape(-1, 1)
work_values = work_encoder.fit_transform(work_resaped)
```

```
print(df_categorical['Geography'][:5])
print()
print(work_values.toarray()[:5])
print()
print(work_encoder.inverse_transform(work_values)[:5])
```

```
0    France
1     Spain
2    France
3    France
4     Spain
```

Name: Geography, dtype: object

```
[[1. 0. 0.]
 [0. 0. 1.]
 [1. 0. 0.]
 [1. 0. 0.]
 [0. 0. 1.]]
```

```
[['France']
 ['Spain']
 ['France']
 ['France']
 ['Spain']]
```

In [154]:

```
df_categorical_encoded = pd.get_dummies(df_categorical, drop_first=True)
df_categorical_encoded.head()
```

Out[154]:

R o w N u m b e r	C u s t o m e r I d	C r e d i t S c o r e	A g e	T e n u r e	B a l a n c e	N u m O f P r o d u c t s	H a s C r C a r d	I s A c t i v e M e m b e r	E s t i m a t e d S a l a r y	S u r n a m e _ Z o t o v a	S u r n a m e _ Z o x	S u r n a m e _ Z u b a r e v	S u r n a m e _ Z u b a r e v	S u r n a m e _ Z u y e v	S u r n a m e _ Z u y e v	S u r n a m e _ Z u y e v	G e o g r a p h y _ G e r m a n y	G e o g r a p h y _ S p a i n	G e n d e r _ M a l e	
		4602			00				.88											
1	2	15647311	608	41	1	83807.86	1	0	1	112542.58	. .42 .58	0	0	0	0	0	0	0	1	0
2	3	15619304	502	42	8	159660.80	3	1	0	113931.57	. .39 .57	0	0	0	0	0	0	0	0	0
3	4	15701354	699	39	1	0.00	2	0	0	93826.63	. .82 .63	0	0	0	0	0	0	0	0	0
4	5	15737888	850	43	2	125510.82	1	1	1	79084.10	. .08 .10	0	0	0	0	0	0	0	1	0

5 rows × 2945 columns

8. Split the data into dependent and independent variables.

Dependent Variable : A dependent variable is a variable whose value depends on another variable.

Independent Variable : An Independent variable is a variable whose value never depends on another variable but the researcher.

In [156]:

```
df=pd.read_csv("E:\IBM projects Assignment Sona  
College\Churn_Modelling.csv")
```

In [157]:

```
print(df["Balance"].min())  
print(df["Balance"].max())  
print(df["Balance"].mean())  
  
0.0  
250898.09  
76485.88928799961
```

In [158]:

```
print(df.count(0))  
  
RowNumber          10000  
CustomerId         10000  
Surname            10000  
CreditScore        10000  
Geography          10000  
Gender             10000  
Age               10000  
Tenure            10000  
Balance           10000  
NumOfProducts     10000  
HasCrCard          10000  
IsActiveMember    10000  
EstimatedSalary   10000  
Exited            10000  
dtype: int64
```

In [159]:

```
print(df.shape)  
  
(10000, 14)
```

In [160]:

```
print(df.size)  
  
140000
```

In [161]:

```
X = df.iloc[:, :-1].values  
print(X)  
  
[[1 15634602 'Hargrave' ... 1 1 101348.88]  
 [2 15647311 'Hill' ... 0 1 112542.58]  
 [3 15619304 'Onio' ... 1 0 113931.57]  
 ...  
 [9998 15584532 'Liu' ... 0 1 42085.58]  
 [9999 15682355 'Sabbatini' ... 1 0 92888.52]  
 [10000 15628319 'Walker' ... 1 0 38190.78]]
```

In [162]:

```
Y = df.iloc[:, -1].values
print(Y)
[1 0 1 ... 1 1 0]
```

9. Scale the independent variables

In [164]:


```
df = pd.read_csv("E:\IBM projects Assignment Sona
College\Churn_Modelling.csv")

x = df[['Age', 'Tenure']].values
y = df['Gender'].values

fig, ax = plt.subplots(ncols=2, figsize=(12, 4))

ax[0].scatter(x[:,0], y)
ax[1].scatter(x[:,1], y)

plt.show()
```




In [165]:

```
fig, ax = plt.subplots(figsize=(12, 4))

ax.scatter(x[:,0], y)
ax.scatter(x[:,1], y)
```

Out[165]:




In [166]:

```
fig, ax = plt.subplots(figsize=(12, 4))
```

```
ax.hist(x[:,0])
ax.hist(x[:,1])
```

Out[166]:

```
(array([ 413., 1035., 1048., 1009.,  989., 1012.,  967., 1028., 1025.,
        1474.]),
 array([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10.]),
)
```



In [167]:


```
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
fig, ax = plt.subplots(figsize=(12, 4))
```

```
scaler = StandardScaler()
x_std = scaler.fit_transform(x)
```

```
ax.hist(x_std[:,0])
ax.hist(x_std[:,1])
```

Out[167]:

```
(array([ 413., 1035., 1048., 1009., 2001.,    0., 1995.,    0., 1025.,
        1474.]),
 array([-1.73331549, -1.38753759, -1.04175968, -0.69598177, -0.35020386,
        -0.00442596,  0.34135195,  0.68712986,  1.03290776,  1.37868567,
         1.72446358]),
)
```




In [168]:

```
fig, ax = plt.subplots(figsize=(12, 4))
```

```
scaler = StandardScaler()
x_std = scaler.fit_transform(x)
```

```
ax.scatter(x_std[:,0], y)
ax.scatter(x_std[:,1], y)
```

Out[168]:



In [169]:


```
fig, ax = plt.subplots(figsize=(12, 4))

scaler = MinMaxScaler()
x_minmax = scaler.fit_transform(x)

ax.hist(x_minmax[:,0])
ax.hist(x_minmax[:,1])
```

Out[169]:

```
(array([ 413., 1035., 1048., 1009.,  989., 1012.,  967., 1028., 1025.,
        1474.]),
 array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1. ]),
 )
```




In [170]:

```
fig, ax = plt.subplots(figsize=(12, 4))

scaler = MinMaxScaler()
x_minmax = scaler.fit_transform(x)

ax.scatter(x_minmax[:,0], y)
ax.scatter(x_minmax[:,1], y)
```

Out[170]:




In [171]:

```
fig, ax = plt.subplots(figsize=(12, 4))

scaler = MinMaxScaler()
x_minmax = scaler.fit_transform(x)

ax.scatter(x_minmax[:,0], y)
```

Out[171]:



```
fig, ax = plt.subplots(figsize=(12, 4))

scaler = MinMaxScaler() x_minmax = scaler.fit_transform(x)

ax.hist(x_minmax[:,0])
```

In [173]:

```
from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.linear_model import SGDRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_absolute_error
import sklearn.metrics as metrics

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

# Import Data
df = pd.read_csv("E:\IBM projects Assignment Sona
College\Churn_Modelling.csv")
x = df[['Age', 'Tenure']].values
```

```

y = df['Balance'].values

# Split into a training and testing set
X_train, X_test, Y_train, Y_test = train_test_split(x, y)

# Define the pipeline for scaling and model fitting
pipeline = Pipeline([
    ("MinMax Scaling", MinMaxScaler()),
    ("SGD Regression", SGDRegressor())
])

# Scale the data and fit the model
pipeline.fit(X_train, Y_train)

# Evaluate the model
Y_pred = pipeline.predict(X_test)
print('Mean Absolute Error: ', mean_absolute_error(Y_pred, Y_test))
print('Score', pipeline.score(X_test, Y_test))

Mean Absolute Error: 56212.80728015005
Score 0.0015999284466322594

```

10. Split the data into training and testing

In [175]:

```

dataset = pd.read_csv("E:\IBM projects Assignment Sona
College\Churn_Modelling.csv")
print(dataset)

```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
\							
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43
...
9995	9996	15606229	Obijiaku	771	France	Male	39
9996	9997	15569892	Johnstone	516	France	Male	35
9997	9998	15584532	Liu	709	France	Female	36
9998	9999	15682355	Sabbatini	772	Germany	Male	42
9999	10000	15628319	Walker	792	France	Female	28

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	
...	
9995	5	0.00	2	1	0	
9996	10	57369.61	1	1	1	
9997	7	0.00	1	0	1	
9998	3	75075.31	2	1	0	

```
9999          4  130142.79          1          1          0
```

```

EstimatedSalary  Exited
0          101348.88      1
1          112542.58      0
2          113931.57      1
3           93826.63      0
4           79084.10      0
...          ...      ...
9995          96270.64      0
9996         101699.77      0
9997          42085.58      1
9998          92888.52      1
9999          38190.78      0

```

```
[10000 rows x 14 columns]
```

In [176]:

```
dataset.drop(["HasCrCard"],axis=1,inplace=True)
```

In [177]:

```
print(dataset.shape)#no. of rows and colume
print(dataset.head(10))
```

```
(10000, 13)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	
5	6	15574012	Chu	645	Spain	Male	44	
6	7	15592531	Bartlett	822	France	Male	50	
7	8	15656148	Obinna	376	Germany	Female	29	
8	9	15792365	He	501	France	Male	44	
9	10	15592389	H?	684	France	Male	27	

	Tenure	Balance	NumOfProducts	IsActiveMember	EstimatedSalary	Exited
0	2	0.00	1	1	101348.88	
1	1	83807.86	1	1	112542.58	
2	8	159660.80	3	0	113931.57	
3	1	0.00	2	0	93826.63	
4	2	125510.82	1	1	79084.10	
5	8	113755.78	2	0	149756.71	
6	7	0.00	2	1	10062.80	
7	4	115046.74	4	0	119346.88	
8	4	142051.07	2	1	74940.50	
9	2	134603.88	1	1	71725.73	

In [178]:

```
X=dataset.iloc[:, :-1].values  
X
```

Out[178]:

```
array([[1, 15634602, 'Hargrave', ..., 1, 1, 101348.88],  
       [2, 15647311, 'Hill', ..., 1, 1, 112542.58],  
       [3, 15619304, 'Onio', ..., 3, 0, 113931.57],  
       ...,  
       [9998, 15584532, 'Liu', ..., 1, 1, 42085.58],  
       [9999, 15682355, 'Sabbatini', ..., 2, 0, 92888.52],  
       [10000, 15628319, 'Walker', ..., 1, 0, 38190.78]], dtype=object)
```

In [179]:

```
Y=dataset.iloc[:, -1].values  
Y
```

Out[179]:

```
array([1, 0, 1, ..., 1, 1, 0], dtype=int64)
```

In [185]:

```
from sklearn.model_selection import train_test_split  
X_train,X_test,Y_train,Y_test = train_test_split( X, Y, test_size = 0.25,  
random_state = 0 )
```

In [186]:

```
print(X_test)  
[[9395 15615753 'Upchurch' ... 1 1 192852.67]  
 [899 15654700 'Fallaci' ... 1 0 128702.1]  
 [2399 15633877 'Morrison' ... 1 1 75732.25]  
 ...  
 [2042 15709846 'Yeh' ... 1 0 84487.62]  
 [1109 15678886 'Golubev' ... 2 0 46522.68]  
 [3333 15720508 'Hsing' ... 1 0 72927.68]]
```

In []: